

# Project #2: Scheduling

**Instructor: Sungyong Ahn**

# Process Control Block in xv6

## ■ struct proc in proc.h

```
35 enum procstate { UNUSED, EMBRYO, SLEEPING, RUNNABLE, RUNNING, ZOMBIE };
36
37 // Per-process state
38 struct proc {
39     uint sz;                // Size of process memory (bytes)
40     pde_t* pgdir;          // Page table
41     char *kstack;          // Bottom of kernel stack for this process
42     enum procstate state;   // Process state
43     int pid;               // Process ID
44     struct proc *parent;    // Parent process
45     struct trapframe *tf;   // Trap frame for current syscall
46     struct context *context; // switch() here to run process
47     void *chan;            // If non-zero, sleeping on chan
48     int killed;            // If non-zero, have been killed
49     struct file *ofile[NOFILE]; // Open files
50     struct inode *cwd;      // Current directory
51     char name[16];         // Process name (debugging)
52 };
53
54 // Process memory is laid out contiguously, low addresses first:
55 //   text
56 //   original data and bss
57 //   fixed-size stack
58 //   expandable heap
```

# Process Control Block in xv6

## ■ struct ptable in proc.c

```
10 struct {  
11     struct spinlock lock;  
12     struct proc proc[NPROC];  
13 } ptable;
```

## ■ in param.h

```
1 #define NPROC      64 // maximum number of processes
```

## ■ Process states (procstate in proc.h)

- UNUSED: Not used
- EMBRYO: Newly allocated (not ready for running yet)
- SLEEPING: Waiting for I/O, child process, or time
- RUNNABLE: Ready to run
- RUNNING: Running on CPU
- ZOMBIE: Exited

# Xv6 Process Scheduler

## ■ main() in main.c

```
int
main(void)
{
    userinit();           // first user process
    mpmain();             // finish this processor's setup
}
```

## ■ mpmain() in main.c

```
// Common CPU setup code.
static void
mpmain(void)
{
    cprintf("cpu%d: starting %d\n", cpuid(), cpuid());
    idtinit();           // load idt register
    xchg(&(mycpu()->started), 1); // tell startothers() we're up
    scheduler();         // start running processes
}
```

# Xv6 Process Scheduler (Cont'd)

- `scheduler()` in `proc.c`
  - Round-robin fashion

Choose a next  
process

Start to execute  
chosen process

```

322 void
323 scheduler(void)
324 {
325     struct proc *p;
326     struct cpu *c = mycpu();
327     c->proc = 0;
328
329     for(;;){
330         // Enable interrupts on this processor.
331         sti();
332
333         // Loop over process table looking for process to run.
334         acquire(&ptable.lock);
335         for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
336             if(p->state != RUNNABLE)
337                 continue;
338
339             // Switch to chosen process. It is the process's job
340             // to release ptable.lock and then reacquire it
341             // before jumping back to us.
342             c->proc = p;
343             switchvm(p);
344             p->state = RUNNING;
345
346             switch(&(c->scheduler), p->context);
347             switchkvm();
348
349             // Process is done running for now.
350             // It should have changed its p->state before coming back.
351             c->proc = 0;
352         }
353         release(&ptable.lock);
354     }
355 }
356 
```

# Xv6 Process Scheduler (Cont'd)

## ■ swtch () in swtch.S

Read **Ch.5 Scheduling**  
of xv6 commentary book!!

in proc.h

```
struct context {
    uint edi;
    uint esi;
    uint ebx;
    uint ebp;
    uint eip;
};
```

```

9  .globl swtch
10 swtch:
11     movl 4(%esp), %eax
12     movl 8(%esp), %edx
13
14     # Save old callee-save registers
15     pushl %ebp
16     pushl %ebx
17     pushl %esi
18     pushl %edi
19
20     # Switch stacks
21     movl %esp, (%eax)
22     movl %edx, %esp
23
24     # Load new callee-save registers
25     popl %edi
26     popl %esi
27     popl %ebx
28     popl %ebp
29     ret
```

# Xv6 Entering Scheduler

## ■ When?

1. Exiting process (`exit()` in `proc.c`)

```
264 // Jump into the scheduler, never to return.
265 curproc->state = ZOMBIE;
266 sched();
```

2. Sleeping process (`sleep()` in `proc.c`)

```
438 // Go to sleep.
439 p->chan = chan;
440 p->state = SLEEPING;
441
442 sched();
```



# Xv6 Entering Scheduler (Cont'd)

## ■ When?

### 3. Yielding CPU due to timer interrupt

- `trap()` in `trap.c`

```
103 // Force process to give up CPU on clock tick.
104 // If interrupts were on while locks held, would need to check nlock.
105 if(myproc() && myproc()->state == RUNNING &&
106     tf->trapno == T_IRQ0+IRQ_TIMER)
107     yield();
```

- `yield()` in `proc.c`

```
384 // Give up the CPU for one scheduling round.
385 void
386 yield(void)
387 {
388     acquire(&ptable.lock); //DOC: yieldlock
389     myproc()->state = RUNNABLE;
390     sched();
391     release(&ptable.lock);
392 }
```

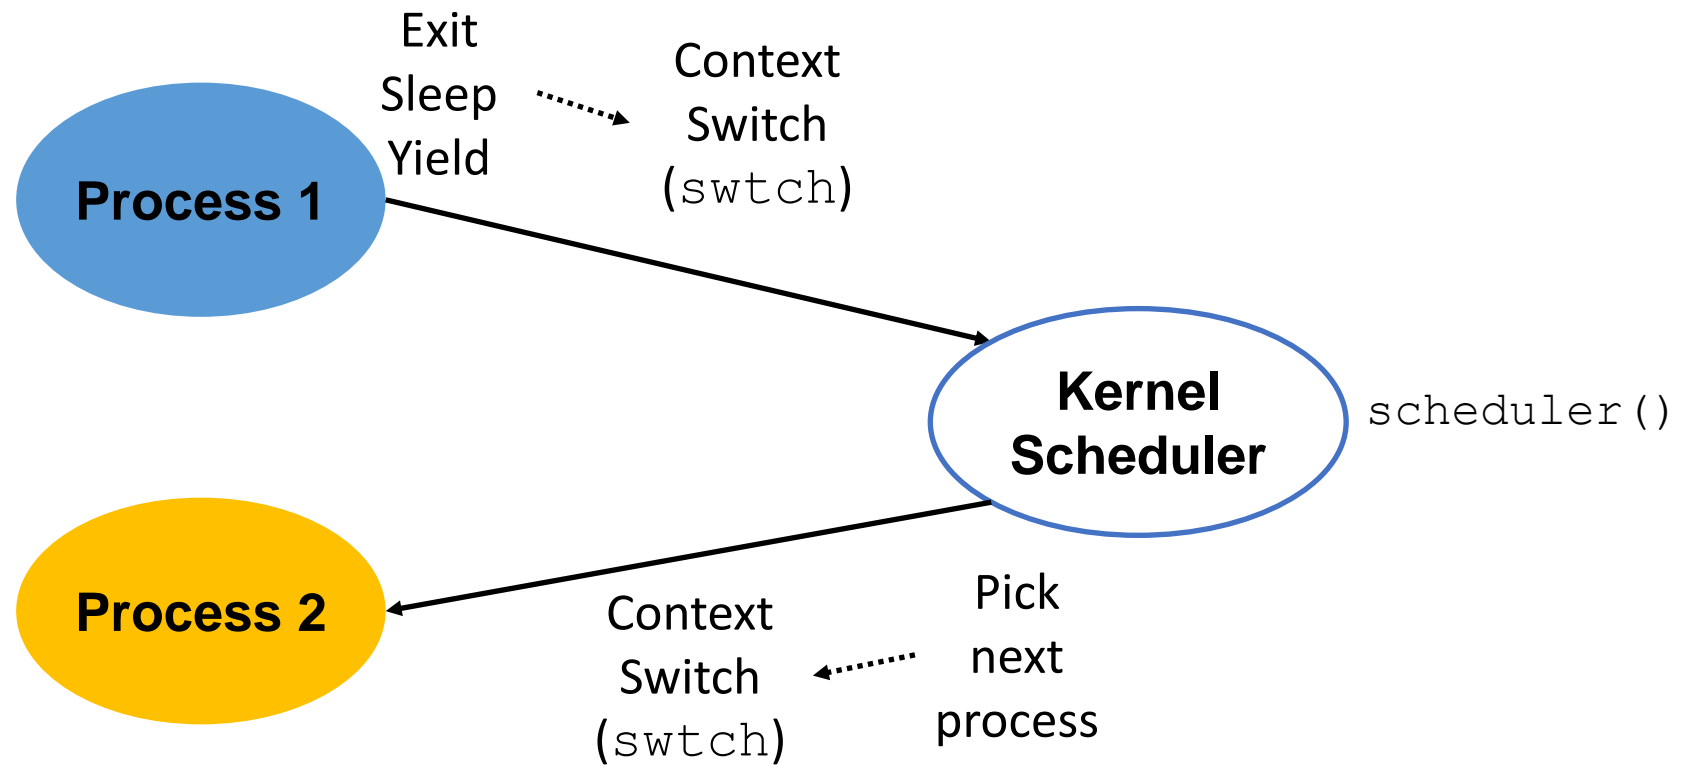


# Xv6 Entering Scheduler (Cont'd)

## ■ sched() in proc.c

```
358 // Enter scheduler. Must hold only ptable.lock
359 // and have changed proc->state. Saves and restores
360 // intena because intena is a property of this
361 // kernel thread, not this CPU. It should
362 // be proc->intena and proc->ncli, but that would
363 // break in the few places where a lock is held but
364 // there's no process.
365 void
366 sched(void)
367 {
368     int intena;
369     struct proc *p = myproc();
370
371     if(!holding(&ptable.lock))
372         panic("sched ptable.lock");
373     if(mycpu()->ncli != 1)
374         panic("sched locks");
375     if(p->state == RUNNING)
376         panic("sched running");
377     if(readeflags() & FL_IF)
378         panic("sched interruptible");
379     intena = mycpu()->intena;
380     swtch(&p->context, mycpu()->scheduler);
381     mycpu()->intena = intena;
382 }
```

# Xv6 Scheduler Summary



# Project #2. Priority Scheduler

## 1. Implement **system calls** related to process priority

- `setnice()` , `getnice()` , `ps()`

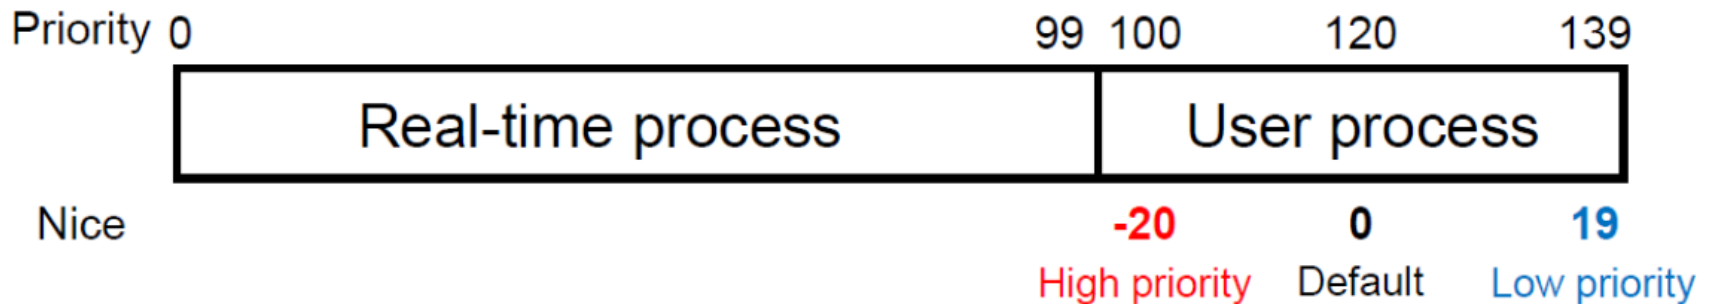
## 2. Implement **priority-based scheduler** on xv6

- The lower nice value, the higher priority
- The highest priority process is selected for next running
  - Tiebreak: FIFO fashion

# Project #2-1: Process Priority

- For represent weight between processes

- In Linux



- In Xv6

- Not implemented!

# Project #2-1: Process Priority

- Make three system calls (`getnice`, `setnice`, `ps`)
- Description of nice value
  - The default nice value is 20.
  - The range of valid nice value is [0, 40].
  - Lower nice values cause more favorable scheduling.
  - When a process calls `fork()` system call, the nice value of child process is same as its parent process.

# Project #2-1: Process Priority

- Make three system calls (`getnice`, `setnice`, `ps`)

- `int setnice(int pid, int nice_value)`

- Get the nice value of process (min: 0 / max: 40)
- Return 0 on success.
- Return -1 if there is no process corresponding to the pid or the nice value is invalid.

- `int getnice(int pid)`

- Return the nice value of pid(process)
- Return -1 if there is no process corresponding to the pid.

# Project #2-1: Process Priority

## ■ `void ps(int pid)`

- The `ps` function prints out process(s)'s information, which includes **pid**, **ppid**, **nice**, **status**, and **name of each process**. If the pid is 0, print out all processes' information. Otherwise, print out corresponding process's information. If there is no process corresponding to the pid, print out nothing.
- No return value
- Make a user program (**minitop**) displays process status using `ps` system call

```
qemu-system-i386 -nographic -drive file=xv6...
cpu1: starting 1
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nl
init: starting sh
$ minitop
pid      ppid      prio      state     name
1         1         20       sleep    init
2         1         20       sleep    sh
3         2         20       run      minitop
$
```

■ Reading **Chap. 3** of `xv6-commentary` will help your project

# Project #2-1: Process Priority

## ■ To-do list for process priority

- Create a variable for a nice value
- Initialization code for a nice value
  - What is default nice value?
  - How can child process inherit the nice value of parent process?

## ■ Then, implement three system calls related to nice value

- `setnice()`, `getnice()`, `ps()`



# fork()

## ■ in proc.c

```
180 int
181 fork(void)
182 {
183     int i, pid;
184     struct proc *np;
185     struct proc *curproc = myproc();
186
187     // Allocate process.
188     if((np = allocproc()) == 0){
189         return -1;
190     }
191
192     // Copy process state from proc.
193     if((np->pgdir = copyvm(curproc->pgdir, curproc->sz)) == 0){
194         kfree(np->kstack);
195         np->kstack = 0;
196         np->state = UNUSED;
197         return -1;
198     }
199     np->sz = curproc->sz;
200     np->parent = curproc;
201     *np->tf = *curproc->tf;
```

# Process Dump in proc.c

```
//PAGEBREAK: 36
// Print a process listing to console.  For debugging.
// Runs when user types ^P on console.
// No lock to avoid wedging a stuck machine further.
void
procdump(void)
{
    static char *states[] = {
        [UNUSED]    "unused",
        [EMBRYO]     "embryo",
        [SLEEPING]   "sleep ",
        [RUNNABLE]   "runble",
        [RUNNING]    "run   ",
        [ZOMBIE]     "zombie"
    };
    int i;
    struct proc *p;
    char *state;
    uint pc[10];

    for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
        if(p->state == UNUSED)
            continue;
        if(p->state >= 0 && p->state < NELEM(states) && states[p->state])
            state = states[p->state];
        else
            state = "???";
        cprintf("%d %s %s", p->pid, state, p->name);
        if(p->state == SLEEPING){
            getcallerpcs((uint*)p->context->ebp+2, pc);
            for(i=0; i<10 && pc[i] != 0; i++)
                cprintf(" %p", pc[i]);
        }
        cprintf("\n");
    }
}
```

# Process Dump in proc.c

```
user@ubuntu:~/xv6-SSE3044$ make qemu-nox
dd if=/dev/zero of=xv6.img count=10000
10000+0 records in
10000+0 records out
5120000 bytes (5.1 MB, 4.9 MiB) copied, 0.01316 s, 389 MB/s
dd if=bootblock of=xv6.img conv=notrunc
1+0 records in
1+0 records out
512 bytes copied, 0.000992285 s, 516 kB/s
dd if=kernel of=xv6.img seek=1 conv=notrunc
334+1 records in
334+1 records out
171164 bytes (171 kB, 167 KiB) copied, 0.000756917 s, 226 MB/s
qemu-system-i386 -nographic -drive file=fs.img,index=1,media=disk,format=raw -drive f
ormat=raw -smp 2 -m 512
xv6...
cpul: starting 1
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
init: starting sh
$ 1 sleep  init 80103dc7 80103e69 80104837 801058d9 80105638
2 sleep  sh 80103d8c 801002c2 80100f8c 80104b32 80104837 801058d9 80105638
$
```

# Project #2-2. Priority Scheduler

## ■ Ignore to yield CPU on clock tick (trap ()) in trap.c)

```
// Force process to give up CPU on clock tick.  
// If interrupts were on while locks held, would need to check nlock.  
/*  
if(myproc && myproc->state == RUNNING && tf->trapno == T_IRQ0+IRQ_TIMER)  
    yield();  
*/
```

## ■ Entering scheduler when

1. Exiting process
2. Sleeping process
3. Yielding CPU
4. Changing priority (setnice)

# Sample Test Program

```
#include "types.h"
#include "stat.h"
#include "user.h"

int main(int argc, char **argv){
    int pid;
    int mypid;

    setnice(1, 19);
    setnice(getpid(), 2);

    pid = fork();

    if(pid == 0){    //Child
        printf(1, "##### State 2 #####\n");
    }
    else{           //Parent
        setnice(pid, 10);    //Set nice value of Child
        printf(1, "##### State 1 #####\n");
        wait();              //Scheduling
        printf(1, "##### State 3 #####\n");
    }

    mypid = getpid();
    printf(1, "PID %d is finished\n", mypid);

    exit();
}
```

```
$ test2
##### State 1 #####
##### State 2 #####
PID 4 is finished
##### State 3 #####
PID 3 is finished
```

# Project #2. Template Code

- Download **xv6-pnu-p2.tar.gz** from PLATO
- Modifications
  - CPUS=1 (Makefile)

```
ifndef CPUS  
CPUS := 1  
endif
```

- Ignore to yield CPU on clock tick
- **yield** system call
  - Yield CPU

# Submission

## ■ Compress your xv6 folder as `StudentID-2.tar.gz`

- `$make clean`
- `$tar -czvf StudentID-2.tar.gz ./xv6-pnu-p2`
- Please command `$make clean` before compressing

## ■ Submit your `tar.gz` file through **PLATO**

## ■ Due date: **4/20, 23:59**

- -25% penalty of total mark per day

## ■ **PLEASE DO NOT COPY !!**

- YOU WILL GET F IF YOU COPIED