

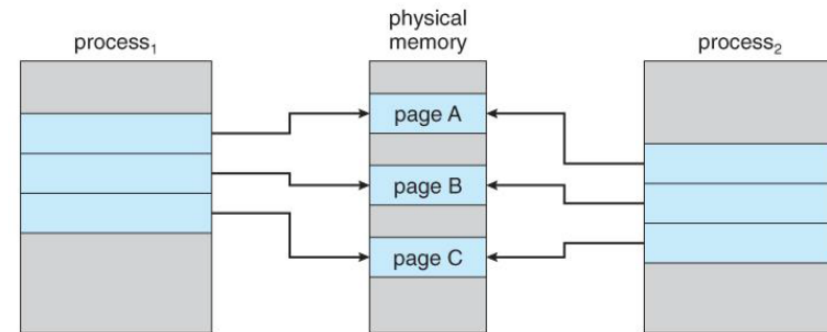
Project #4: Copy-On-Write

Instructor: Sungyong Ahn

Copy-on-Write

■ When a process forks

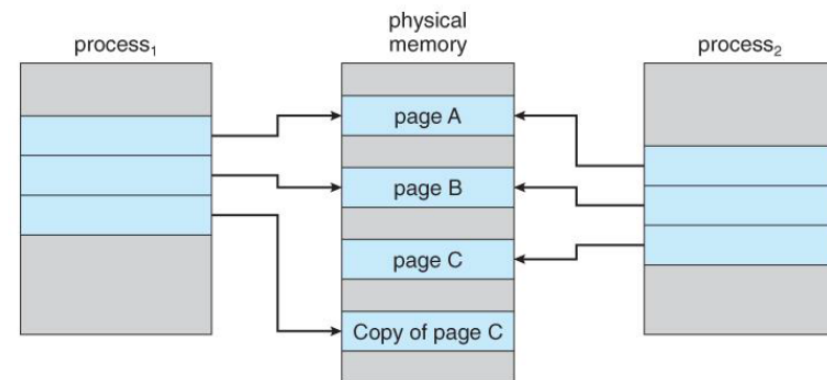
- Create shared mappings to the same page frames in physical page
- Shared pages are protected as **read-only**



Before process 1 modifies page C

■ When data is written to shared pages

- Protection fault is generated
- OS allocates new space in physical memory and directs the write to it



After process 1 modifies page C

■ Reference counter for physical pages is needed

Project #3. Copy-on-Write

■ proc.c

```
int
fork(void)
{
    int i, pid;
    struct proc *np;
    struct proc *curproc = myproc();

    // Allocate process.
    if((np = allocproc()) == 0){
        return -1;
    }

    // Copy process state from proc.
    if((np->pgdir = copyuvm(curproc->pgdir, curproc->sz)) == 0){
        kfree(np->kstack);
        np->kstack = 0;
        np->state = UNUSED;
        return -1;
    }
}
```

vm.c

copyuvm()

```

pde_t*
copyuvm(pde_t *pgdir, uint sz)
{
    pde_t *d;
    pte_t *pte;
    uint pa, i, flags;
    char *mem;

    if((d = setupkvm()) == 0)
        return 0;
    for(i = 0; i < sz; i += PGSIZE){
        if((pte = walkpgdir(pgdir, (void *) i, 0)) == 0)
            panic("copyuvm: pte should exist");
        if(!(*pte & PTE_P))
            panic("copyuvm: page not present");
        pa = PTE_ADDR(*pte);
        flags = PTE_FLAGS(*pte);
        if((mem = kalloc()) == 0)
            goto bad;
        memmove(mem, (char*) P2V(pa), PGSIZE);
        if(mappages(d, (void*) i, PGSIZE, V2P(mem), flags) < 0)
            goto bad;
    }
    return d;

bad:
    freevm(d);
    return 0;
}

```

kernel address space를 위한 page table 생성

pde_t *pgdir은 parent process의 page table
pde_t *d는 새로운 child process를 위한 page table

vm . C

copyvm ()

```
pde_t*
copyvm(pde_t *pgdir, uint sz)
{
```

```
    pde_t *d;
    pte_t *pte;
    uint pa, i, flags;
    char *mem;
```

pde_t *pgdir은 parent process의 page table
pde_t *d는 새로운 child process를 위한 page table

```
    if((d = setupkvm()) == 0)
        return 0;
```

```
    for(i = 0; i < sz; i += PGSIZE){
        if((pte = walkpgdir(pgdir, (void *) i, 0)) == 0)
```

virtual address에 대한 page
table entry 반환

```
        panic("copyvm: pte should exist");
```

```
        if(!(*pte & PTE_P))
```

```
            panic("copyvm: page not present");
```

```
        pa = PTE_ADDR(*pte);
```

```
        flags = PTE_FLAGS(*pte);
```

```
        if((mem = kalloc()) == 0)
```

```
            goto bad;
```

```
        memmove(mem, (char*) P2V(pa), PGSIZE);
```

```
        if(mappages(d, (void*)i, PGSIZE, V2P(mem), flags)
```

```
            goto bad;
```

page table entry에서 user virtual
address에 대한 physical
address를 구함

주의: pa는 parent process의
physical page address

```
    }
    return d;
```

```
bad:
    freevm(d);
    return 0;
```

```
}
```

vm . C

copyvm ()

```
pde_t*
copyvm(pde_t *pgdir, uint sz)
{
```

```
    pde_t *d;
    pte_t *pte;
    uint pa, i, flags;
    char *mem;
```

pde_t *pgdir은 parent process의 page table
pde_t *d는 새로운 child process를 위한 page table

```
    if((d = setupkvm()) == 0)
        return 0;
```

```
    for(i = 0; i < sz; i += PGSIZE){
        if((pte = walkpgdir(pgdir, (void *) i, 0)) == 0)
            panic("copyvm: pte should exist");
        if(!(*pte & PTE_P))
            panic("copyvm: page not present");
```

새로운 physical page 할당
mem은 새로 할당 받은 page의
virtual address

```
        pa = PTE_ADDR(*pte);
        flags = PTE_FLAGS(*pte);
        if((mem = kalloc()) == 0)
            goto bad;
```

새로운 physical page에
parent의 physical page를 복사

```
        memmove(mem, (char*) P2V(pa), PGSIZE);
```

```
        if(mappages(d, (void*) i, PGSIZE, V2P(mem), flags) < 0)
            goto bad;
```

```
    }
    return d;
```

```
bad:
    freevm(d);
    return 0;
```

}

vm.c

copyuvm()

```
pde_t*
copyuvm(pde_t *pgdir, uint sz)
{
```

```
    pde_t *d;
    pte_t *pte;
    uint pa, i, flags;
    char *mem;
```

pde_t *pgdir은 parent process의 page table
pde_t *d는 새로운 child process를 위한 page table

```
    if((d = setupkvm()) == 0)
        return 0;
    for(i = 0; i < sz; i += PGSIZE){
        if((pte = walkpgdir(pgdir, (void *) i, 0)) == 0)
            panic("copyuvm: pte should exist");
        if(!(*pte & PTE_P))
            panic("copyuvm: page not present");
        pa = PTE_ADDR(*pte);
        flags = PTE_FLAGS(*pte);
        if((mem = kalloc()) == 0)
            goto bad;
        memmove(mem, (char*) P2V(pa), PGSIZE);
        if(mappages(d, (void*)i, PGSIZE, V2P(mem), flags) < 0)
            goto bad;
    }
    return d;

bad:
    freevm(d);
    return 0;
}
```

virtual address에 대한 page table
entry를 child process의 page
table에 생성

Project #4. Copy-on-Write

- Implement copy-on-write on xv6
- Implementation details
 1. Modify `copyvm()` in `vm.c` from copy version to duplicate version

Modify copyvm()

- 새로운 page 할당없이 parent process의 page를 child process의 mapping table에 매핑
- page table entry(PTE)에 writeable flag (PTE_W) disable
 - 해당 page에 write를 시도하면 page fault 발생 -> Copy-on-write 수행
- physical page의 reference count 증가
- 1cr3 (V2P (pgdir)) 를 호출해 TLB flush

```
for(i = 0; i < sz; i += PGSIZE){
    if((pte = walkpgdir(pgdir, (void *) i, 0)) == 0)
        panic("copyvm: pte should exist");
    if(!(*pte & PTE_P))
        panic("copyvm: page not present");
    pa = PTE_ADDR(*pte);
    flags = PTE_FLAGS(*pte);
    if((mem = kalloc()) == 0)
        goto bad;
    memmove(mem, (char*) P2V(pa), PGSIZE);
    if(mappages(d, (void*)i, PGSIZE, V2P(mem), flags) < 0)
        goto bad;
}
```

Project #4. Copy-on-Write

- Implement copy-on-write on xv6
- Implementation details
 1. Modify `copyvm()` in `vm.c` from copy version to duplicate version
 2. Managing page reference counter

Reference Counter for Physical Pages

■ kalloc.c

- 각 physical page의 reference counter를 기록할 배열 생성

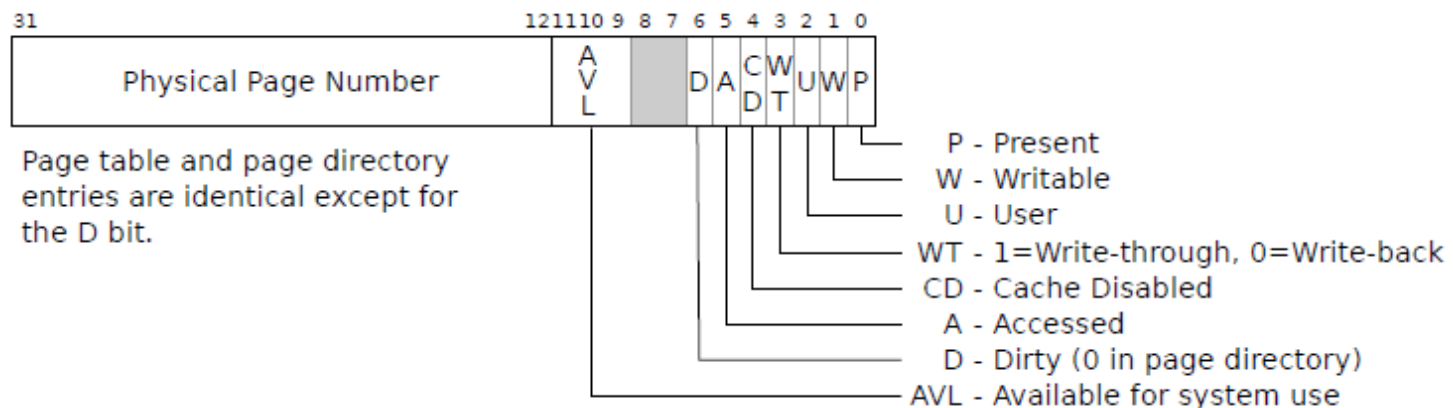
- uint `pgrefcount` [`PHYSTOP >> PGSHIFT`];

- `memlayout.h`

```
#define PHYSTOP 0xE000000 // Top physical memory
```

- `mmu.h`

```
#define PGSHIFT 12 // log2(PGSIZE)
```



Reference Counter for Physical Pages

■ kalloc.c

- void freerange(void *vstart, void *vend)
 - xv6 commentary book page 33
 - Reference counter 초기화

```
void
freerange(void *vstart, void *vend)
{
    char *p;
    p = (char*) PGROUNDUP((uint)vstart);
    for(; p + PGSIZE <= (char*)vend; p += PGSIZE)
        kfree(p);
}
```

Reference Counter for Physical Pages

■ kalloc.c

- void kfree(char *v)
 - reference counter 감소
 - reference counter가 zero가 되면 해당 page를 free page list에 추가

```
void
kfree(char *v)
{
    struct run *r;

    if((uint)v % PGSIZE || v < end || V2P(v) >= PHYSTOP)
        panic("kfree");

    // Fill with junk to catch dangling refs.
    memset(v, 1, PGSIZE);

    if(kmem.use_lock)
        acquire(&kmem.lock);
    numfreepages++;
    r = (struct run*)v;
    r->next = kmem.freelist;
    kmem.freelist = r;
    if(kmem.use_lock)
        release(&kmem.lock);
}
```

Reference Counter for Physical Pages

■ kalloc.c

- char* kalloc(void)
 - reference counter 1로 초기화

```
char*
kalloc(void)
{
    struct run *r;

    if(kmem.use_lock)
        acquire(&kmem.lock);
    numfreepages--;
    r = kmem.freelist;
    if(r)
        kmem.freelist = r->next;
    if(kmem.use_lock)
        release(&kmem.lock);
    return (char*)r;
}
```

Reference Counter for Physical Pages

■ `kalloc.c`

- Reference counter 관련 API 구현
- `uint get_refcounter(uint pa)`
 - 인자로 받은 physical address가 속한 physical page의 reference counter를 반환
- `void dec_refcounter(uint pa)`
 - physical page의 reference counter를 1 감소
- `void inc_refcounter(uint pa)`
 - physical page의 reference counter를 1 증가

Project #4. Copy-on-Write

- Implement copy-on-write on xv6
- Implementation details
 1. Modify `copyvm()` in `vm.c` from copy version to duplicate version
 2. Managing page reference counter
 3. Implementing page-fault handler

Page Fault Handler Implementation

```
■ trap.c void
trap(struct trapframe *tf)
{
    if(tf->trapno == T_SYSCALL) {
        if(myproc()->killed)
            exit();
        myproc()->tf = tf;
        syscall();
        if(myproc()->killed)
            exit();
        return;
    }

    switch(tf->trapno) {
        case T_PGFLT:
            pagefault();
            break;
        case T_IRQ0 + IRQ_TIMER:
            if(cpuid() == 0) {
                acquire(&tickslock);
                ticks++;
                wakeup(&ticks);
                release(&tickslock);
            }
            lapiceoi();
            break;
    }
}
```

Page Fault Handler Implementation

■ vm.c

- void pagefault(void)
 - rcr2() 를 호출해 page fault가 발생한 virtual address 결정
 - virtual address가 valid한 값인지 확인
 - 해당 virtual address의 page table entry(pte_t) 찾기
 - walkpgdir()
 - physical address와 reference counter 찾기
 - get_refcounter()
 - reference counter가 1보다 큰 경우
 - 새로운 페이지를 할당 받아 기존 페이지를 복사
 - » kalloc(), memmove()
 - page table entry 내용 변경
 - reference counter 1 감소 (dec_refcounter())
 - reference counter가 1인 경우
 - 현재 page table entry에 writeable flag만 enable
 - lcr3(V2P(pgdir)) 를 호출해 TLB flush

Test programs

■ test1.c

```
int main(int argc, char **argv)
{
    int before, after;
    int pid;

    printf(1, "TEST1: ");
    before = freemem();
    pid = fork();
    if(pid == 0){
        after = freemem();
        if(before - after == 68)
            printf(1, "OK\n");
        else
            printf(1, "WRONG\n");
        exit();
    }
    else{
        wait();
    }
    exit();
}
```

```
$ test1
TEST1: OK
$ test2
TEST2: OK
$ test3
TEST3: OK
$ test4
TEST4: OK
```

Project #4. Template Code

- Download **xv6-pnu-4.tar.gz** from PLATO
- Modifications
 - `freemem()` system call
 - Return the number of free pages in `kmem.freelist`

Submission

- **Compress your xv6 folder as `StudentID-4.tar.gz`**

- `$tar -czvf StudentID-4.tar.gz ./xv6-pnu-4`
- Please command `$make clean` before compressing

- **Submit your `tar.gz` file through **PLATO****

- **Due date: **6/10 (Thur.), 23:59****

- **Late submission penalty**

- -25% penalty of total mark per day

- **PLEASE DO NOT COPY !!**

- YOU WILL GET F IF YOU COPIED

Tips

- Reading xv6 commentary will help you a lot
 - <https://pdos.csail.mit.edu/6.828/2017/xv6/book-rev10.pdf>
 - The line numbers in this book refer to the source booklet below
 - Reading **chap. 2 “Page tables”** of xv6-commentary will help your project
 - <https://pdos.csail.mit.edu/6.828/2017/xv6/xv6-rev10.pdf>
- 프로젝트 관련 질문은 PLATO 질의응답 게시판을 활용